

# Homework 5

## Problem 1

### 1.a Gibbs sampler

```
gibbs_sampler <- function(X, Y, alpha_init = rep(0,dim(X)[2]), sigma2_init = 1,
                          delta = 10, lambda = 10, mu = rep(0,dim(X)[2]),
                          Omega = diag(dim(X)[2]), MCsamplesize = 1000){

  # dims of X
  n <- dim(X)[1] # sample size
  p <- dim(X)[2] # number of covariates in linear regression model/alpha-parameters

  # set starting values for parameter chains
  # variables where generated values will be saved
  alpha_chain <- matrix(NA, MCsamplesize, p)
  sigma2_chain <- rep(NA, MCsamplesize)
  # first values for both chains
  alpha_chain[1,] <- alpha_init # [0] * p
  sigma2_chain[1] <- sigma2_init # [1]

  # create parameter chain
  for (i in 2:(MCsamplesize)){

    alpha_chain[i,] <- alpha_chain[i-1,]
    sigma2_chain[i] <- sigma2_chain[i-1]

    for (j in 1:p){

      s_j <- 1 / (Omega[j, j] + sigma2_chain[i]^(-1) * sum(X[, j]^2))
      m_j <- s_j*(mu[j]*Omega[j, j]
                  - sum(Omega[j, -j]*(alpha_chain[i, -j] - mu[-j]))
                  + sigma2_chain[i]^(-1)*sum(X[, j]*(Y - X[, -j]*alpha_chain[i, -j])))
      alpha_chain[i,j] <- rnorm(1, m_j, sqrt(s_j)) # rnorm accepts sd, not var

    }

    sigma2_chain[i] <- 1 / rgamma(n=1, shape=delta + n/2,
                                rate=lambda + 1/2*sum( (Y - X*alpha_chain[i, ])^2 ))
  }

  return(list("alpha" = alpha_chain, "sigma2" = sigma2_chain))
}
```

## 1.b

```
set.seed(999)
samplesize <- 200
dim         <- 3
alpha_real <- c(-2, 0, 4)
sigma_real <- 2
X <- matrix(rnorm(samplesize*dim), nrow = samplesize, n = dim)
Y <- as.vector(X%*%alpha_real + rnorm(samplesize, 0, sigma_real)) # sd, not var

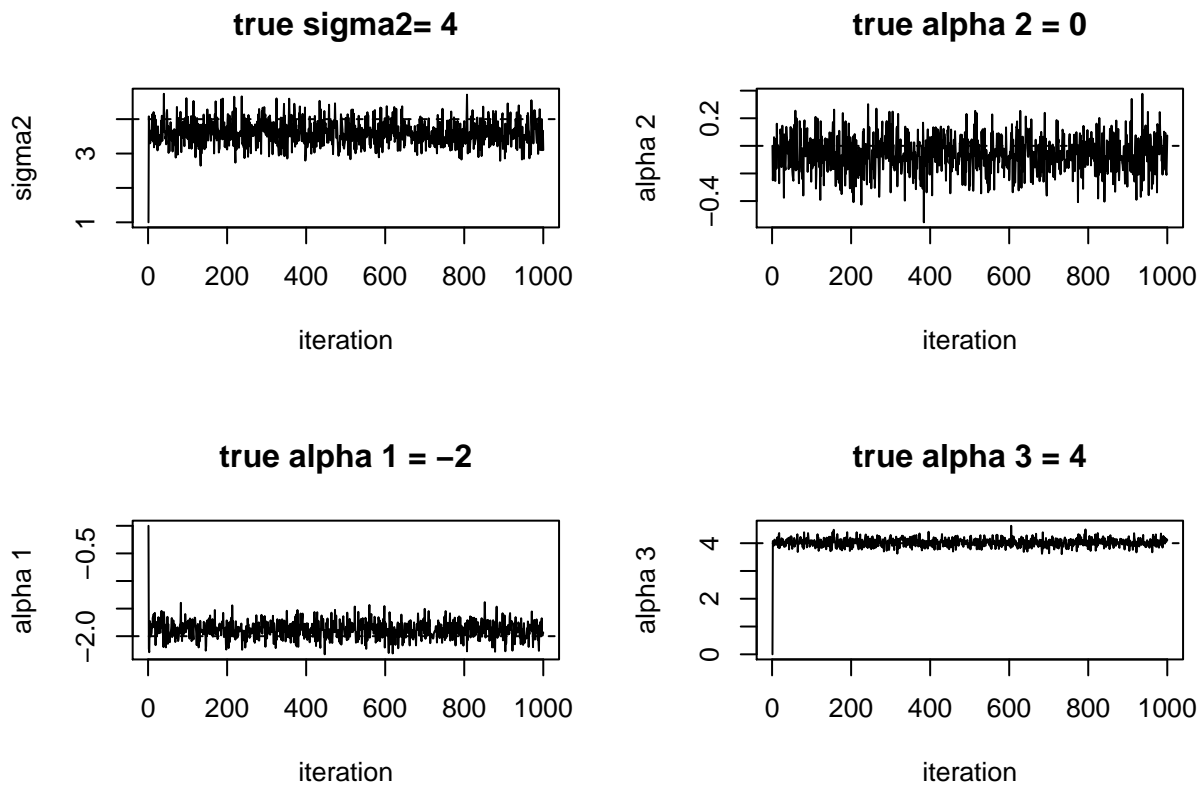
n <- dim(X)[1] # sample size
p <- dim(X)[2] # number of covariates in linear regression model/alpha-parameters

sigma2_real <- sigma_real^2

out <- gibbs_sampler(X, Y)
sigma2 <- out$sigma2
alpha <- out$alpha

par(mfcol = c(2, 2))
plot(sigma2, ylab = "sigma2", xlab = "iteration", main = paste("true sigma2=",
  round(sigma2_real, digits = 2)), type = "l")
abline(h = sigma2_real, lty = 2)

for (i in 1:p){
  plot(alpha[,i], ylab = paste("alpha",i), xlab = "iteration", main = paste(
    "true alpha", i, "=", round(alpha_real[i], digits = 2)), type = "l")
  abline(h = alpha_real[i], lty = 2)
}
```

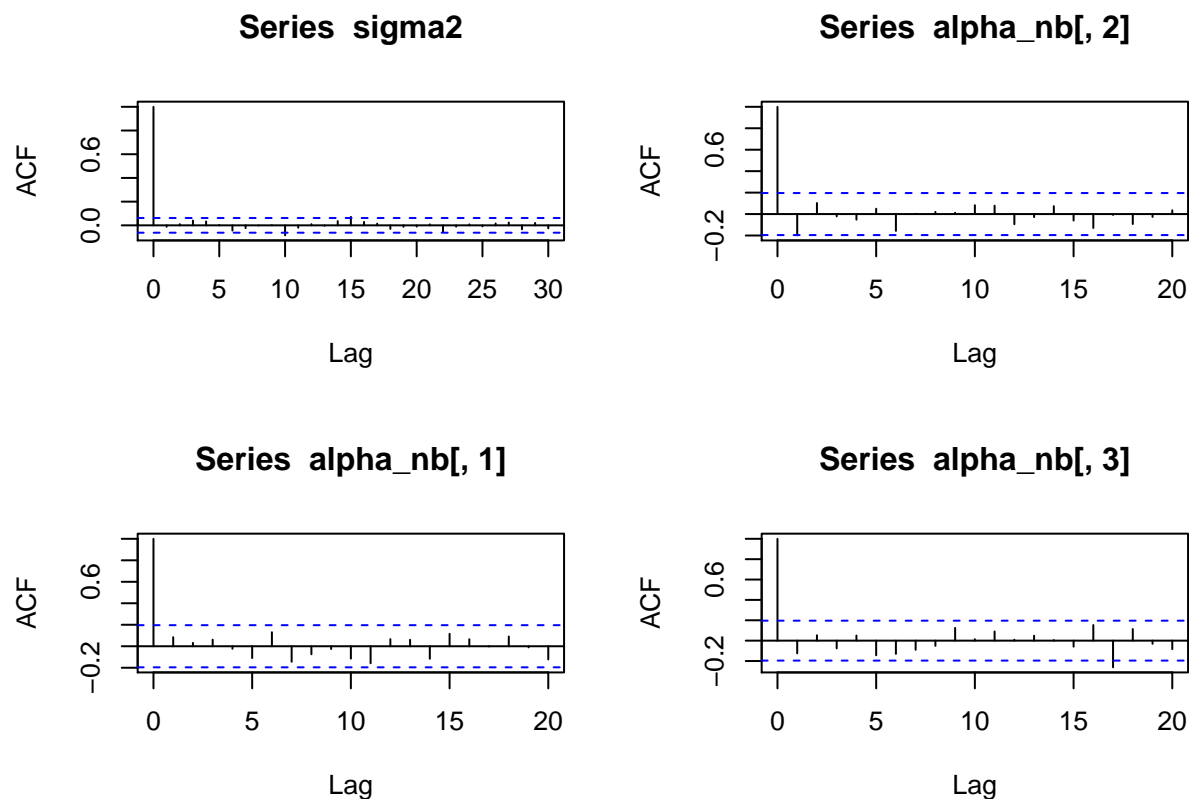


### Intrepret traceplots

As we can see, the chains converge quite fast. From the graph it can be seen that the chains converge in 10-20 iterations. Let's choose a burn-in period of 100 to be safe.

```
# no burnin values
sigma2_nb <- sigma2[101:n]
alpha_nb <- alpha[101:n,]
```

```
par(mfcol = c(2, 2))
acf(sigma2)
acf(alpha_nb[,1])
acf(alpha_nb[,2])
acf(alpha_nb[,3])
```



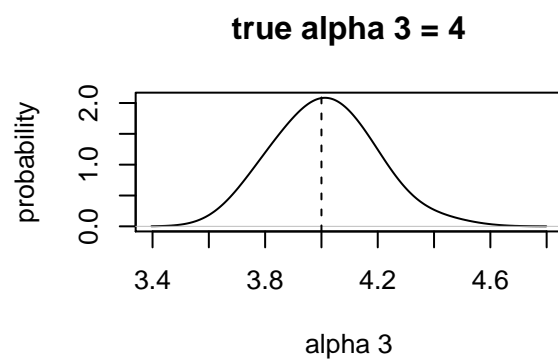
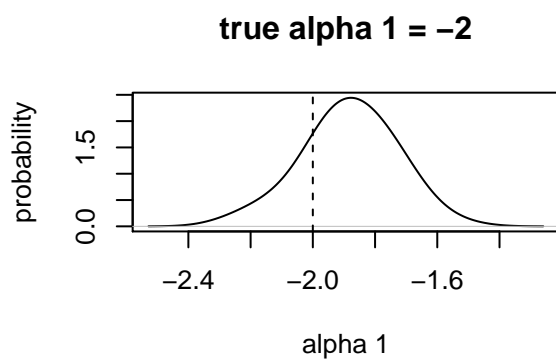
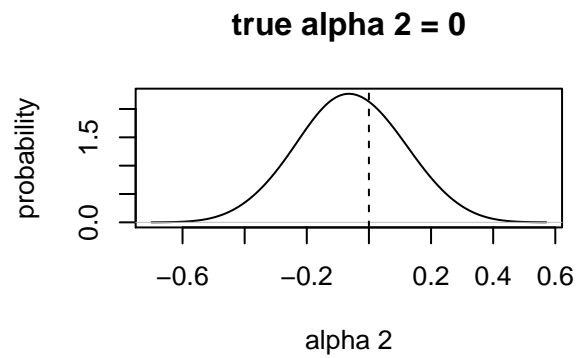
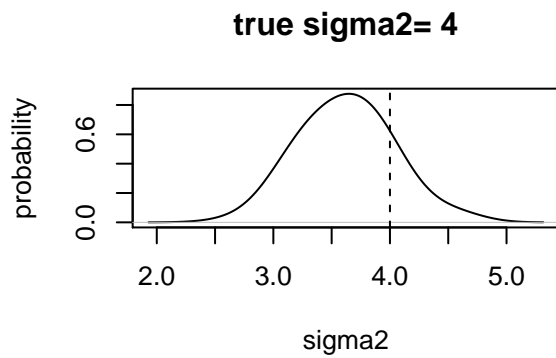
### Intrepret ACF plots

We can see, that autocorrelation is low for each of the chains. So we can proceed without throwing away every n'th sample.

```
# density estimates
par(mfcol = c(2, 2))

hw <- 2 * (summary(sigma2_nb)[5] - summary(sigma2_nb)[2])
plot(density(sigma2_nb, width = hw), xlab = "sigma2", ylab = "probability",
     type = "l", main = "")
abline(v = sigma2_real, lty = 2)
title(paste("true sigma2=", round(sigma2_real, digits = 2)))

for (i in 1:p){
  hw <- 2 * (summary(alpha_nb[,i])[5] - summary(alpha_nb[,i])[2])
  plot(density(alpha_nb[,i], width = hw), xlab = paste("alpha", i), ylab = "probability",
       type = "l", main = "")
  abline(v = alpha_real[i], lty = 2)
  title(paste("true alpha", i, "=", round(alpha_real[i], digits = 2)))
}
```



```
ord_lr <- lm(Y ~ X - 1) # -1 to get read of intercept
ord_lr
```

```
##
## Call:
## lm(formula = Y ~ X - 1)
##
## Coefficients:
##      X1      X2      X3
## -1.9238 -0.0587  4.1003
```

```
summary(ord_lr)
```

```
##
## Call:
## lm(formula = Y ~ X - 1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.3286 -1.5298 -0.1412  1.1600  5.2405
##
## Coefficients:
##      Estimate Std. Error t value Pr(>|t|)
## X1  -1.9238     0.1510 -12.740  <2e-16 ***
## X2  -0.0587     0.1363  -0.431   0.667
```

```
## X3    4.1003    0.1465  27.980   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.956 on 197 degrees of freedom
## Multiple R-squared:  0.8355, Adjusted R-squared:  0.833
## F-statistic: 333.4 on 3 and 197 DF,  p-value: < 2.2e-16
```

## Intrepret lm results

Coefficients of the linear regression are -1.9238 -0.0587 4.1003 and estimated standard error is 1.956, which means that estimated  $\sigma^2=3.83$ . These values are closer to the true values, than the maximum of the posterior densities estimated by Gibbs sampler.

Especially bad is the posterior of  $\sigma^2$  with peak around 3.5.

## 1.c

```
set.seed(42)
samplesize <- 200
dim         <- 3
alpha_real <- c(-2, 0, 4)
sigma_real <- 2
X <- 0.1*matrix(rnorm(samplesize*dim), nrow = samplesize, ncol = dim) +
  matrix(rep(rnorm(samplesize), 3), nrow = samplesize, ncol = dim)
Y <- as.vector(X%%alpha_real + rnorm(samplesize, 0, sigma_real))
```

```
alpha_init <- rep(20,3)
```

```
sigma_init <- 1
tau_init <- 1 / sigma_init^2
```

```
mah_model <- function(){
  for (i in 1:n){
    Y.mu[i] <- X[i,] %*% alpha
    Y[i] ~ dnorm(Y.mu[i], tau)
  }

  tau ~ dgamma(delta, lambda)
  alpha ~ dmnorm(mu, Omega)

  sigma2 <- 1 / tau
}

mah_data=list(X=X, Y=Y, n=samplesize,
             delta = 10,
             lambda = 10,
             mu = rep(0,dim(X)[2]),
             Omega = diag(dim(X)[2]))
```

```

mah_mcmc <- R2jags::jags(
  data = mah_data,
  inits=list(list(alpha=alpha_init, tau=tau_init)),
  parameters.to.save = c("alpha","tau", "sigma2"),
  n.iter = samplesize,
  n.chains = 1,
  n.thin = 1,
  n.burnin = 0, # 100 iterations are run for adaption
  model.file = mah_model)

```

```
## module glm loaded
```

```

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 200
##   Unobserved stochastic nodes: 2
##   Total graph size: 1221
##
## Initializing model

```

```
structure(mah_mcmc)
```

```

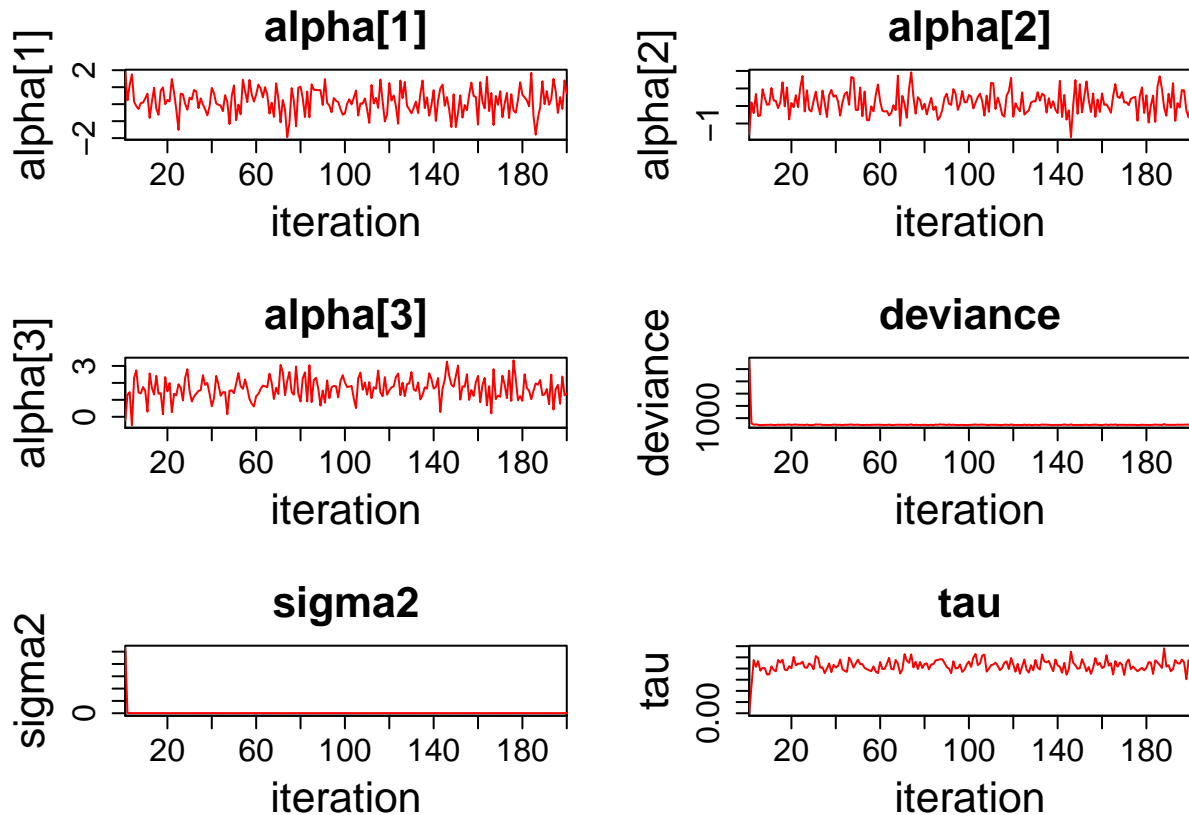
## Inference for Bugs model at "/var/folders/7n/fcgysl1141n_czqxkwtdd68r0000gn/T//RtmpV3zX0F/model114dfef
## 1 chains, each with 200 iterations (first 0 discarded)
## n.sims = 200 iterations saved
##           mu.vect sd.vect   2.5%    25%    50%    75%   97.5%
## alpha[1]  0.180   0.748  -1.310 -0.301  0.168  0.666  1.481
## alpha[2]  0.212   0.661  -0.853 -0.277  0.193  0.630  1.609
## alpha[3]  1.665   0.668   0.266  1.268  1.659  2.122  2.970
## sigma2    17.903 186.095   3.805  4.387  4.704  5.067  5.772
## tau        0.212   0.028   0.173  0.197  0.213  0.228  0.263
## deviance 896.589  74.490 886.429 889.195 890.731 893.061 898.248
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 2774.4 and DIC = 3671.0
## DIC is an estimate of expected predictive error (lower deviance is better).

```

```

rafalib::bigpar(3,2)
R2jags::traceplot(mah_mcmc,mfrow=c(3,2))

```



mah\_mcmc

```
## Inference for Bugs model at "/var/folders/7n/fcgysll141n_czqxkwtdd68r0000gn/T//RtmpV3zX0F/model14dfel"
## 1 chains, each with 200 iterations (first 0 discarded)
## n.sims = 200 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%   75%  97.5%
## alpha[1]  0.180  0.748 -1.310 -0.301  0.168  0.666  1.481
## alpha[2]  0.212  0.661 -0.853 -0.277  0.193  0.630  1.609
## alpha[3]  1.665  0.668  0.266  1.268  1.659  2.122  2.970
## sigma2    17.903 186.095  3.805  4.387  4.704  5.067  5.772
## tau        0.212  0.028  0.173  0.197  0.213  0.228  0.263
## deviance 896.589  74.490 886.429 889.195 890.731 893.061 898.248
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 2774.4 and DIC = 3671.0
## DIC is an estimate of expected predictive error (lower deviance is better).
```



## Problem 2

### 2.a

#### Full Conditionals

$$\begin{aligned}
\pi(\mu_1, \mu_2) &\sim \ell(\mu_1, \mu_2) \cdot p(\mu_1, \mu_2) \\
&= \ell(\mu_1, \mu_2) \cdot 1 \\
&= \ell(\mu_1, \mu_2) \\
&= N_n(y_1, \dots, y_n \mid \mu, \Sigma) \\
&\sim \prod_{i=1}^n \exp \left\{ -\frac{1}{2} (y_i - \mu)^\top \Sigma^{-1} (y_i - \mu) \right\} \\
&\sim \prod_{i=1}^n \exp \left\{ -\frac{1}{2} [y_i^\top \Sigma^{-1} y_i - 2\mu^\top \Sigma^{-1} y_i + \mu^\top \Sigma^{-1} \mu] \right\} \\
&\sim \prod_{i=1}^n \exp \left\{ -\frac{1}{2} [-2\mu^\top \Sigma^{-1} y_i + \mu^\top \Sigma^{-1} \mu] \right\} \\
&= \exp \left\{ -\frac{1}{2} \sum_{i=1}^n [-2\mu^\top \Sigma^{-1} y_i + \mu^\top \Sigma^{-1} \mu] \right\} \\
&= \exp \left\{ -\frac{n}{2} [-2\mu^\top \Sigma^{-1} \bar{y} + \mu^\top \Sigma^{-1} \mu] \right\} \\
&\sim \exp \left\{ -\frac{n}{2} [(\mu - \bar{y})^\top \Sigma^{-1} (\mu - \bar{y})] \right\} \\
&\sim N_2(\mu \mid \bar{y}, \frac{\Sigma}{n}) \\
\pi(\mu_1, \mu_2) &\sim N_2(\mu \mid \bar{y}, \frac{\Sigma}{n}) \\
&\text{where } \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i
\end{aligned}$$

With formula 1.5:

$$\begin{aligned}
\pi(\mu_1 \mid \mu_2) &\sim N(\bar{y}_1 + \frac{\rho\sigma_1}{\sigma_2}(\mu_2 - \bar{y}_2), \frac{\sigma_2^2(1-\rho^2)}{n}) \\
\pi(\mu_2 \mid \mu_1) &\sim N(\bar{y}_2 + \frac{\rho\sigma_2}{\sigma_1}(\mu_1 - \bar{y}_1), \frac{\sigma_1^2(1-\rho^2)}{n}) \mid \\
\bar{y}_1 &= \sum_{i=1}^n y_i^{(1)} \\
\bar{y}_2 &= \sum_{i=1}^n y_i^{(2)}
\end{aligned}$$

These are the full conditionals

#### Gibbs sampling algorithm

With  $\pi(\mu_1 \mid \mu_2)$  and  $\pi(\mu_2 \mid \mu_1)$  as defined above:

1. Set the iteration counter to  $j=1$  and set initial values  $\mu^{(0)} = (\mu_1^{(0)}, \mu_2^{(0)})'$
2. Obtain a new value  $\mu^{(j)} = (\mu_1^{(j)}, \mu_2^{(j)})'$  through successive generation of values

$$\begin{aligned}
\mu_1^{(j)} &\sim \pi(\mu_1 \mid \mu_2^{(j-1)}) \\
\mu_2^{(j)} &\sim \pi(\mu_2 \mid \mu_1^{(j)})
\end{aligned}$$

3. Change counter  $j$  to  $j+1$  and return to step 2 until convergence is reached.

## 2.b

```
gibbs_sampler_2 <- function(Y, sigma1, sigma2, rho,
                           R=1000){
  n <- dim(Y)[1]

  # variables where generated values will be saved
  mu_chain <- matrix(NA, R, 2)
  mu_chain[1,] <- runif(2,0,10)

  # same for all iterations
  sd1 <- sqrt(sigma2^2*(1-rho^2) / n)
  sd2 <- sqrt(sigma1^2*(1-rho^2) / n)
  y.mean1 <- mean(Y[,1])
  y.mean2 <- mean(Y[,2])

  # create parameter chain
  for (i in 2:(R)){
    mu_chain[i, 1] = rnorm(1,
                          mean=(y.mean1 + rho*sigma1/sigma2 * (mu_chain[i-1,2] - y.mean2)),
                          sd=sd1)
    mu_chain[i, 2] = rnorm(1,
                          mean=(y.mean2 + rho*sigma2/sigma2 * (mu_chain[i,1] - y.mean1)),
                          sd=sd2)
  }

  return(mu_chain)
}
```

```
set.seed(42)
mu_real = c(6,2)
sigma1 = sqrt(2)
sigma2 = sqrt(0.5)
rho = 0.05
n = 100
Sig = matrix(c(sigma1^2, sigma1*sigma2*rho,
               sigma1*sigma2*rho, sigma2^2), nrow = 2, ncol = 2)

Y <- MASS::mvrnorm(n, mu_real, Sig)
```

```
R1 = 100
R2 = 100000
```

```
mu_chain_1 <- gibbs_sampler_2(Y, sigma1, sigma2, rho, R=R1)
mu_chain_2 <- gibbs_sampler_2(Y, sigma1, sigma2, rho, R=R2)
```

```
compute_D <- function(mu, Y, sigma1, sigma2, rho){
  y.mean1 <- mean(Y[,1])
  y.mean2 <- mean(Y[,2])
  result <- 1/(1-rho^2) * (
    (mu[1] - y.mean1)^2 / sigma1^2 +
    (mu[2] - y.mean2)^2 / sigma2^2 +
```

```

    2*rho*(mu[1] - y.mean1)*(mu[2] - y.mean2) / (sigma1*sigma2)
  )
  return(result)
}

```

```

DIC_4.12 <- function(mu_chain, Y, sigma1, sigma2, rho){

  R <- dim(mu_chain)[1]

  mu.mean = c(mean(mu_chain[,1]), mean(mu_chain[,2]))
  D_of_mu.mean <- compute_D(mu.mean, Y, sigma1, sigma2, rho)

  D.mean = 0
  for (i in 1:R){
    D_of_mu.i <- compute_D(mu_chain[i,], Y, sigma1, sigma2, rho)
    D.mean <- D.mean + D_of_mu.i
  }
  D.mean <- 1/R * D.mean

  return(D.mean - D_of_mu.mean)
}

```

```

DIC_4.13 <- function(mu_chain, Y, sigma1, sigma2, rho){

  R <- dim(mu_chain)[1]

  D.mean = 0
  for (i in 1:R){
    D_of_mu.i <- compute_D(mu_chain[i,], Y, sigma1, sigma2, rho)
    D.mean <- D.mean + D_of_mu.i
  }
  D.mean <- 1/R * D.mean

  D.var = 0
  for (i in 1:R){
    D_of_mu.i <- compute_D(mu_chain[i,], Y, sigma1, sigma2, rho)
    D.var <- D.var + (D_of_mu.i - D.mean)^2
  }
  D.var <- 1/(2*(R-1)) * D.var

  return(D.var)
}

```

```
DIC_4.12(mu_chain_1, Y, sigma1, sigma2, rho)
```

```
## [1] 0.3692384
```

```
DIC_4.12(mu_chain_2, Y, sigma1, sigma2, rho)
```

```
## [1] 0.04280324
```

```
DIC_4.13(mu_chain_1, Y, sigma1, sigma2, rho)
```

```
## [1] 5.630098
```

```
DIC_4.13(mu_chain_2, Y, sigma1, sigma2, rho)
```

```
## [1] 0.001665639
```

For both DIC from eq. 4.12 and eq. 4.13, MCMC with more iterations obtained a far lower score, which indicates that longer sampling is beneficial.